

Problem 1

For any $w \in \Sigma^*$, let $P_w(x)$ denote the polynomial obtained by feeding w into the automaton. By definition, this function is a polynomial in one variable x . Let X be the set defined in the problem statement.

We consider two cases:

1. There exists a word $w \in \Sigma^*$ such that $P_w(x)$ is not the zero polynomial.

In this case, let $R(P_w)$ represent the set of roots of $P_w(x)$. Since $P_w(x)$ is not identically zero, $R(P_w)$ is finite. Furthermore, X must be a subset of $R(P_w)$; otherwise, there would exist some $x \in X$ for which $P_w(x) \neq 0$, contradicting the definition of X . Thus, X is finite.

2. For every word $w \in \Sigma^*$, $P_w(x)$ is the zero polynomial.

In this scenario, $P_w(x) = 0$ holds for all $x \in \mathbb{Q}$ and every $w \in \Sigma^*$. Consequently, $X = \mathbb{Q}$.

From these cases, we conclude that X is either finite or equal to \mathbb{Q} , completing the proof.

Problem 2

This problem was solved in collaboration with Kacper Bal and Mateusz Mroczka.

Let \mathcal{A} denote the automaton from the problem statement, with dimension d and alphabet Σ . For any $a \in \Sigma$, let $M_a \in \mathbb{Q}^{d \times d}$ represent the matrix corresponding to the transition function for symbol a .

We will construct a polynomial automaton \mathcal{B} that defines a function $g : \Sigma^* \rightarrow \mathbb{Q}$. The goal is for this function to satisfy the property that, for any word $w \in \Sigma^*$, $g(w) = f(w_{\text{lex}})$, where w_{lex} is w sorted lexicographically, and f is the function computed by \mathcal{A} .

Assume $\Sigma = \{a_1, a_2, \dots, a_k\}$, where $a_1 <_{\text{lex}} a_2 <_{\text{lex}} \dots <_{\text{lex}} a_k$. Denote the initial state of \mathcal{A} as $q_{\text{ini}} \in \mathbb{Q}^{1 \times d}$, and the final mapping as $q_{\text{fin}} \in \mathbb{Q}^{d \times 1}$. According to the definition of a weighted automaton, for any $w \in \Sigma^*$ such that $w = w_1 w_2 \dots w_n$, it holds that $f(w) = q_{\text{ini}} \cdot M_{w_1} \cdot M_{w_2} \cdot \dots \cdot M_{w_n} \cdot q_{\text{fin}}$. Therefore, it should also hold that

$$g(w) = q_{\text{ini}} \cdot M_{a_1}^{\#_{a_1}(w)} \cdot M_{a_2}^{\#_{a_2}(w)} \cdot \dots \cdot M_{a_k}^{\#_{a_k}(w)} \cdot q_{\text{fin}}.$$

The automaton \mathcal{B} will have dimension $d^2 \cdot k$. The first d^2 coordinates will correspond to M_{a_1} raised to the power of the number of letters a_1 read so far, the next d^2 coordinates to M_{a_2} , and so on. In the initial state of \mathcal{B} , we store the identity matrix I for each letter a_i in the alphabet. For any $a \in \Sigma$, the transition matrix simulates matrix multiplication by M_a on the appropriate coordinates, while the other coordinates remain unchanged. Note that this requires only linear mappings.

Once the values $M_{a_1}^{\#_{a_1}(w)}, M_{a_2}^{\#_{a_2}(w)}, \dots, M_{a_k}^{\#_{a_k}(w)}$ have been computed and stored in the state, we need to calculate the final value

$$q_{\text{ini}} \cdot M_{a_1}^{\#_{a_1}(w)} \cdot M_{a_2}^{\#_{a_2}(w)} \cdot \dots \cdot M_{a_k}^{\#_{a_k}(w)} \cdot q_{\text{fin}}.$$

This is a polynomial involving $d^2 \cdot k$ variables, which are the elements of the matrices stored in the state. Hence, the automaton \mathcal{B} can compute this, making it the only non-linear transition.

Since \mathcal{A} is also a polynomial automaton (as linear functions are polynomials), we can use the algorithm presented in the lecture to determine whether \mathcal{A} and \mathcal{B} are equivalent. If they are, the function f is commutative; otherwise, it is not, which concludes the proof.